

OBSERVABILITY (18%)

The recommend time for this section is **21,6 minutes** or less.

1. Create an nginx pod with a liveness probe that just runs the command 'ls'. Save its YAML in pod.yaml. Run it, check its probe status, delete it.
2. Modify the pod.yaml file so that liveness probe starts kicking in after 5 seconds whereas the interval between probes would be 5 seconds. Run it, check the probe, delete it.
3. Create an nginx pod (that includes port 80) with an HTTP readinessProbe on path '/' on port 80. Again, run it, check the readinessProbe, delete it.
4. Create a busybox pod that runs 'i=0; while true; do echo "\$i: \$(date)"; i=\$((i+1)); sleep 1; done'. Check its logs
5. Debugging
Create a busybox pod that runs 'ls /notexist'. Determine if there's an error (of course there is), see it. In the end, delete the pod
6. Create a busybox pod that runs 'notexist'. Determine if there's an error (of course there is), see it. In the end, delete the pod forcefully with a 0 grace period
7. Get CPU/memory utilization for nodes (metrics-server must be running)